



Department of Artificial Intelligence and Machine Learning

Date: 06.06.2025	Test - 2	Max. Marks: 50+10
Semester: VI	UG	Duration: 1 $\frac{1}{2}$ Hrs + $\frac{1}{2}$ Hr
Course Title: Big Data Technologies		Course Code: AI362IA

Sl No	Questions	M	BT	CO
1.	Tasks are considered failed if they exceed a timeout period of _____ minutes and it can set by _____ property	2	2	1
2.	The maximum number of attempts given to run a Map Reduce job on application master is _____ and it can be controlled by _____ property	2	2	1
3.	Blacklisting of a Node Manager is performed by the _____, and tasks are rescheduled on a different Node Manager if more than _____ tasks fail	2	2	1
4.	List two differences between Traditional Databases and Hive?	2	2	1
5.	State the two index types used in HiveQL?	2	2	1
PART- B				
1	Summarize the Job Submission and Job Initialization stages in the anatomy of a Map Reduce job run with a diagram.	10	2	1
2(a)	Explain how progress and status updates are handled during a MapReduce job run and identify their main components.	3+3	2	1
2(b)	Explain how a Node Manager failure occurs in Hadoop and how it is handled by the system.	4	2	2
3(a)	Outline the Shuffle and Sort phase from the perspective of the Reduce side in a Map Reduce job.	6	2	2
3(b)	State the reasons for disabling speculative execution in a MapReduce job.	4	2	2
4(a)	Discuss how hive organises tables into partitions?	5	2	2
4(b)	Explain the sorting and aggregating concepts in hive with an example?	5	2	2
5	<p>A music shop records information about its customers, including their customer ID, first name, and last name. The shop also keeps track of all the movies available, identified by a movie ID and movie name.</p> <p>Orders placed by the shop are tracked using an order ID, along with the supplier ID, movie ID, and the number of copies ordered.</p> <p>Details of all suppliers who provide the CDs are maintained, including their supplier ID and supplier name.</p> <p>Finally, the renting information is recorded with the customer ID, tape ID, checkout date, and rental duration.</p> <p>For the above scenario, perform the following</p> <ol style="list-style-type: none"> Create relevant tables Load the data into the table Write following queries using HiveQL <ol style="list-style-type: none"> List the movies that have more than 4 copies been ordered? Which customers rented movies for 5 days or more? Which movies aren't in the inventory? Which movie was rented for the longest duration (by any customer)? 	10	3	4